

Image Processing With Python

Stephen Granade

@Sargent

stephen@granades.com

WTF is Image Processing?

- Modify images
 - Change colors
 - Adjust brightness and contrast
 - Remove parts of an image
- Extract information (signal) from images
 - Identify what's in an image (like faces)
 - Track objects as they move

Five Things

- Data structure to hold the image data
- Way to read in image data
- Ways to manipulate the image data



Three ~~Five~~ Things

- Data structure to hold the image data
- Way to read in image data
- Ways to manipulate the image data

Three ~~Five~~ Things

- Data structure to hold the image data
 - NumPy arrays
- Way to read in image data
 - SciPy library
- Ways to manipulate the image data
 - SciPy and scikit-image libraries

Numpy, SciPy, Scikits, Matplotlib



+



Get and Use the Libraries

- NumPy, SciPy and Matplotlib: <http://www.scipy.org/>
 - `import numpy as np`
 - `from scipy import ndimage`
 - `from scipy import misc`
 - `import matplotlib.pyplot as plt`
- Scikit-image: <http://scikit-image.org/>
 - `import skimage`
 - Lots of sub-libraries

NumPy Arrays: Ordered Collections of Homogeneous Data

- *Ordered*: You get ahold of items by their index number
- *Collection*: Holds other types of data, like ints or strings or lists or even other arrays
- *Homogeneous*: only holds data of the same type

Arrays are Sequences, But Slices Aren't Copies!

Lists

```
l1 = range(10)
l2 = l1[1:-1]
print('l1: %s\nl2: %s' % (l1, l2))
```

```
l1[4] = -1
print('l1: %s\nl2: %s' % (l1, l2))
```

Arrays

```
a1 = np.arange(10)
a2 = a1[1:-1]
print('a1: %s\na2: %s' % (a1, a2))
```

```
a1[4] = -1
print('a1: %s\na2: %s' % (a1, a2))
```

Arrays are Sequences, But Slices Aren't Copies!

Lists

```
l1 = range(10)
l2 = l1[1:-1]
print('l1: %s\nl2: %s' % (l1, l2))
    l1: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
    l2: [1, 2, 3, 4, 5, 6, 7, 8]
l1[4] = -1
print('l1: %s\nl2: %s' % (l1, l2))
    l1: [0, 1, 2, 3, -1, 5, 6, 7, 8, 9]
    l2: [1, 2, 3, 4, 5, 6, 7, 8]
```

Arrays

```
a1 = np.arange(10)
a2 = a1[1:-1]
print('a1: %s\na2: %s' % (a1, a2))
    a1: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
    a2: [1, 2, 3, 4, 5, 6, 7, 8]
a1[4] = -1
print('a1: %s\na2: %s' % (a1, a2))
    a1: [0, 1, 2, 3, -1, 5, 6, 7, 8, 9]
    a2: [1, 2, 3, -1, 5, 6, 7, 8]
```

Need a copy? Use `a1.copy()`

List Slices are New Objects But Array Slices Aren't

```
l1 = range(10)  
l2 = l1[1:-1]
```

[0 1, 2, 3, 4, 5, 6, 7, 8, 9]
11

[1 2, 3, 4, 5, 6, 7, 8]
12

```
a1 = np.arange(10)  
a2 = a1[1:-1]
```

[0 1, 2, 3, 4, 5, 6, 7, 8, 9]
11


12

Useful Array Functions

```
A = np.ones([3, 2])
```

Function	Result	Notes
<code>np.shape(A)</code> OR <code>A.shape</code>	(3, 2)	Tuple of array dimensions
<code>np.ndim(A)</code> OR <code>A.ndim</code>	2	Number of array dimensions
<code>np.size(A)</code> OR <code>A.size</code>	6	Flat length of the array (number of elements)
<code>A.T</code> OR <code>A.transpose()</code>	2x3 array	Transpose
<code>np.append(A, 2)</code>	7x1 vector	Flatten and append values
<code>np.append(A, [2,2])</code>	8x1 vector	Flatten and append values
<code>np.append(A, [[2,2]], axis=0)</code>	4x2 array	Append values along <i>axis</i> dimension
<code>np.concatenate((A, A))</code>	6x2 array	Given a tuple of arrays, concatenates them

Images Are (Usually) a 2D or 3D Array



Images Are (Usually) a 2D or 3D Array

- 2D array
 - Effectively grayscale image
 - Usually integers, but the same image functions that work on a uint8 grayscale image often work on 2D floating-point or Boolean arrays
- 3D array
 - 2D picture with a third dimension for things like color
 - But how is color defined?

Color Spaces Map Colors to Numbers

- RGB
 - Probably the one you've seen before
 - Red, green, and blue channels (values along the 3rd dimension)
 - Additive color mixing
- CMYK
 - Cyan, magenta, yellow, black
 - Subtractive color mixing
 - Mainly print – doesn't show up much in computer images

Color Spaces Map Colors to Numbers

- YCbCr
 - Luminance, Chroma blue, Chroma red
 - Shows up in JPEG & MPEG
- HSV
 - Hue, Saturation, Value
 - Hue: what color. Saturation: how colorful. Value: kind of how light it is, but not really.
 - Cylindrical coordinate representation of RGB



```
#####  
# SIMPLE IMAGE LOADING AND DISPLAYING  
  
import numpy as np  
from scipy import ndimage  
from scipy import misc  
import matplotlib.pyplot as plt  
  
img = misc.imread('Schroedinger.jpg')  
  
plt.imshow(img)  
  
# Why's it all rainbow like that? Because SCIENCE. We can fix that, though  
plt.close('all')  
plt.imshow(img, cmap='gray')
```

Read In Images Using SciPy

- A lot of the `scipy.misc` and `scipy.ndimage` functions are using the Python Imaging Library behind the scenes, and the SciPy docs about them are sketchy
- `scipy.misc.imread()` figures out an image file's type by the file contents
- `scipy.misc.imsave()` uses the filename extension to determine the file type
 - Supported types: bmp, gif, jpg (or jpeg), png, tiff, xbm
 - For a full list, see <http://effbot.org/imagingbook/formats.htm>

Convert an Image to Python Imaging Library Format for All PIL Options

- `scipy.misc.imsave()` won't let you set options like JPEG quality, but PIL will
- Solution: convert your ndarray image to a PIL image and then use its save function
 - ```
pil_img = misc.toimage(img)
pil_img.save('filename.jpg',
 quality=30)
```
- <http://effbot.org/imagingbook/formats.htm> has information on all save options
- Note: there's both `misc.toimage()` and `misc.fromimage()` for moving between NumPy and PIL

```
#####
CROPPING, SCALING, ROTATING, FLIPPING

Crop using standard array slicing notation
imgy, imgx = img.shape
cropped_img = img[0:imgy // 2, :]

To scale, use the imresize() function from scipy.misc
resized_img = misc.imresize(img, 0.30)
plt.close('all')
compare_images([img, resized_img], title='Float-Resized Image',
 subtitles=['Original', 'Resized'])
imresize takes the size as a float (fraction), integer (percent), or
tuple (final size)...or so it says. As far as I can tell, integer
scaling is broken
resized_img = misc.imresize(img, 10)
Tuples work, though
resized_img = misc.imresize(img, (img.shape[0] // 2, img.shape[1]))
plt.close('all')
compare_images([img, resized_img], title='Tuple-Resized Image')

To rotate, use the rotate() function from scipy.ndimage
rotated_img = ndimage.rotate(img, 30)
plt.close('all')
compare_images([img, rotated_img], title='Rotated Image')
By default rotate will make the array big enough that nothing gets cut
off. You can change that, of course
cropped_rotated_img = ndimage.rotate(img, 30, reshape=False)

To flip, use the standard NumPy functions flipud() and fliplr()
flipped_img = np.flipud(img)
```



1



2



3



4

# Local Filtering Determines a Pixel's Values By Looking At Its Neighbors

- The basic idea: replace a pixel's value by a function involving neighbors' values
- Top right: a 3x3 Gaussian blurring filter. Each pixel's value is partially replaced by the value of its neighbors
- Bottom right: Sobel horizontal filter for finding horizontal lines
- Local filters let you blur, sharpen, detect edges, and more

|        |       |        |
|--------|-------|--------|
| $1/16$ | $1/8$ | $1/16$ |
| $1/8$  | $1/4$ | $1/8$  |
| $1/16$ | $1/8$ | $1/16$ |

|        |        |        |
|--------|--------|--------|
| $1/8$  | $1/4$  | $1/8$  |
| 0      | 0      | 0      |
| $-1/8$ | $-1/4$ | $-1/8$ |



1




2

```
scipy.ndimage includes several common filters. For example, Gaussian
filters, which soften and blur images, are in scipy.ndimage
blurred_img = ndimage.gaussian_filter(img, sigma=1)

The larger the Gaussian's sigma, the more it blurs
more_blurred_img = ndimage.gaussian_filter(img, sigma=3)
```

1

2



```
What if you have a noisy image?
cropped_img = img[50:140, 90:180]
noisy_img = cropped_img + (cropped_img.std()*np.random.random(cropped_img.shape) -
 (cropped_img.std()/2)*np.ones(cropped_img.shape))

You can use a Gaussian filter to de-noise the image
denoised_img = ndimage.gaussian_filter(noisy_img, sigma=1)


Or you can use a median filter to better preserve edges
median_denoised_img = ndimage.median_filter(noisy_img, 3)
plt.close('all')
compare_images([noisy_img, denoised_img, median_denoised_img],
 title="Gaussian vs Median Denoising",
 subtitles=['Noisy', 'Gaussian', 'Median'])
```

# SciKit-Image Adds to SciPy

```
from skimage import...
```

| Subpackage   | Description                                                       |
|--------------|-------------------------------------------------------------------|
| color        | Color space conversion (RGB, YCbCr, etc.)                         |
| data         | Test images and example data                                      |
| draw         | Drawing primitives – lines, text, etc. – for NumPy arrays         |
| exposure     | Image intensity adjustment, like histogram equalization           |
| feature      | Feature detection and extraction                                  |
| filter       | Sharpening, edge-finding, rank filters, thresholding              |
| graph        | Graphi-theoretic operations like finding the shortest path        |
| io           | Reading, saving, and displaying images                            |
| measure      | Measurement of image properties like similarity and contours      |
| morphology   | Morphological operations like opening or skeletonization          |
| restoration  | Restoration algorithms like denoising or deconvolution algorithms |
| segmentation | Partitioning an image into multiple regions                       |
| transform    | Geometric and other transforms like rotation, Radon transform     |






```

LOCAL FILTERING



skimage has a lot of local filters available, like Sobel
import skimage.filter

img = misc.imread('Schroedinger.jpg')
schro_vsobel = skimage.filter.vsobel(img)
```



```
Remember my noise reduction example earlier? skimage has better routines
cropped_img = img[50:140, 90:180]
noisy_img = cropped_img + (
 (cropped_img.std()*np.random.random(cropped_img.shape) -
 (cropped_img.std()/2)*np.ones(cropped_img.shape)))
median_denoised_img = ndimage.median_filter(noisy_img, 3)
total_var_denoised_img = skimage.filter.denoise_tv_chambolle(noisy_img,
 weight=30)
```

# Histogram Equalization Improves Contrast



```
#####
ADJUSTING EXPOSURE

import skimage.exposure



Using skimage, we can perform contrast enhancement automatically by
equalizing the picture's histogram. The presentation has more information
on histogram equalization

sudoku = ndimage.imread('sudoku.jpg', flatten=True)
"flatten=True" reads the picture in as grayscale 1 also converts the
image to floating-point numbers.

In a lot of cases, skimage wants image data to be floating point numbers
scaled to the range [-1, 1].
Because of the flattening, sudoku is a floating-point number, but it's
scaled to the range [0, 255]. Adjust it to a [-1, 1] scale.
sudoku_scaled = (sudoku - 127.5)/256
sudoku_equalized = skimage.exposure.equalize_hist(sudoku_scaled) 2
```



# Matching Images





# Example: Extract the Square Glyphs



---

# How to draw an Owl.

---

*"A fun and creative guide for beginners"*

---



Fig 1. Draw two circles




Fig 2. Draw the rest of the damn Owl

# Example: Extract the Square Glyphs






# The Glyphs are 11x11 B&W Squares





Luminance  
(Brightness)



Chrominance  
(Color)


```
FILENAME = 'fez-image.jpg'

glyphimg = misc.imread(FILENAME)

Read in the image in YCbCr (luminance/chroma blue/chroma red) color space
glyphimg_ycbcr = ndimage.imread(FILENAME, mode='YCbCr')
lum = glyphimg_ycbcr[:, :, 0]
chrom = glyphimg_ycbcr[:, :, 1] / 2 + glyphimg_ycbcr[:, :, 2] / 2
```



Luminance  
(Brightness)




Chrominance  
(Color)

```
FILENAME = 'fez-image.jpg'

glyphimg = misc.imread(FILENAME)


Read in the image in YCbCr (luminance/chroma blue/chroma red) color space
glyphimg_ycbcr = ndimage.imread(FILENAME, mode='YCbCr')
lum = glyphimg_ycbcr[:, :, 0]
chrom = glyphimg_ycbcr[:, :, 1] / 2 + glyphimg_ycbcr[:, :, 2] / 2
```




```
The glyph graphics are an 11x11 square of just black and white surrounded
by a black square. That means we're looking for areas with both high luminance
(brightness) and low chromaticity (color).

lum_mask = lum > .94*np.max(lum)
Low color corresponds to a chrominance of around 128 (halfway between 0 & 255)
chrom_mask = np.logical_and(chrom >= 126, chrom <= 136)

Let's identify the part of the image we care about
image_mask = np.logical_and(lum_mask, chrom_mask)
```




Possible Glyph Pixels  
(Bright, No Color)



Possible Border Pixels  
(Dark, No Color)

```
Find the border around the glyphs (low luminance, low chromaticity)
border_mask = np.logical_and(lum < 20, chrom_mask)
```



```
Routines to find all of the top left corners in an image
def mask_is(mask, i, j):
 """Return the true/false value of a Boolean mask image at index [i, j].

 If the [i, j] index is out of bounds, returns True.
 """
 try:
 return mask[i, j]
 except IndexError:
 return True


def is_top_left_corner(image, border, i, j):
 """Determine if image[i, j] is a top left corner that we're looking for.

 Args

 image : array
 A Boolean array indicating where the image pixels are.
 border : array
 A Boolean array indicating where the border around the image is.

 Returns

 True if image[i, j] is a top left corner, False otherwise.
 """
 return (image[i, j] &
 mask_is(image, i, j+1) &
 mask_is(image, i+1, j+1) &
 mask_is(image, i+1, j) &
 mask_is(border, i+1, j-1) &
 mask_is(border, i, j-1) &
 mask_is(border, i-1, j-1) &
 mask_is(border, i-1, j) &
 mask_is(border, i-1, j+1) &
 (not mask_is(image, i-2, j)) &
 (not mask_is(image, i, j-2)))
```



```
Routines to find all of the top left corners in an image
def mask_is(mask, i, j):
 """Return the true/false value of a Boolean mask image at index [i, j].

 If the [i, j] index is out of bounds, returns True.
 """
 try:
 return mask[i, j]
 except IndexError:
 return True

def is_top_left_corner(image, border, i, j):
 """Determine if image[i, j] is a top left corner that we're looking for.


 Args

 image : array
 A Boolean array indicating where the image pixels are.
 border : array
 A Boolean array indicating where the border around the image is.


 Returns

 True if image[i, j] is a top left corner, False otherwise.
 """
 return (image[i, j] &
 mask_is(image, i, j+1) &
 mask_is(image, i+1, j+1) &
 mask_is(image, i+1, j) &
 mask_is(border, i+1, j-1) &
 mask_is(border, i, j-1) &
 mask_is(border, i-1, j-1) &
 mask_is(border, i-1, j) &
 mask_is(border, i-1, j+1) &
 (not mask_is(image, i-2, j)) &
 (not mask_is(image, i, j-2)))
```

# Convex Hull Stretches a Rubber Band Over Objects







```
Find the likely columns by counting occurrences and find the most-common
spacing between adjacent columns
(I'm using a Counter: a dict that stores elements as keys and the number of
times that element occurs as the values)
columns = Counter(top_left_corners[1])
avg_column_count = np.mean(columns.values())
likely_columns = sorted([k for k in columns
 if columns[k] > avg_column_count])
col_delta, _ = Counter([i - j for i, j
 in zip(likely_columns[1:], likely_columns[:-1])]).most_common(1)[0]
I happen to know that the spacing between rows is 1 more than that between
columns, & there aren't enough columns to make the statistical analysis
I did above work for rows.
(I also know that the column spacing is the glyph size (11 pixels) + 2,
but I wanted to show how I'd do this kind of analysis)
row_delta = col_delta + 1
```

# Where to Go From Here

- Look at the sample code from this presentation
  - <https://github.com/sgranade/python-image-processing-intro>
- Try a more optimized image processing library
  - OpenCV: <http://opencv.org/>
- Try a simpler image processing library
  - SimpleCV: <http://simplecv.org/>
- Look at other tutorials
  - <http://blog.yhathq.com/posts/image-processing-with-scikit-image.html> (scikit-image)
  - <http://www.pyimagesearch.com/> (OpenCV)
  - [http://docs.opencv.org/trunk/doc/py\\_tutorials/py\\_tutorials.html](http://docs.opencv.org/trunk/doc/py_tutorials/py_tutorials.html) (OpenCV)

# Image Processing With Python

Stephen Granade

@Sargent

[stephen@granades.com](mailto:stephen@granades.com)